

In [1]:

```
'''
Dependencies
'''

import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
```

In [2]:

```
'''
Reading the MNIST/USPS Handwritten Digits Dataset
'''

def readData(fname='ZipDigits.train'):
    '''
        Input:
            fname: name of file containing N examples, each
with d attributes
        Output:
            X: N x d+1 numpy array
            y: N x 1 numpy array
    '''
    with open(fname) as f:
        X = []
        y = []
        rlines = f.readlines()
        for line in rlines:
            row = line.rstrip().split(' ')
            yval = int(float(row[0]))
            y.append(yval)
            xvals = [float(pixel) for pixel in row[1:]]
            X.append(xvals)
        X = np.array(X)
        y = np.array(y)
        y = y.reshape((y.shape[0], 1))
        print(f'X shape: {X.shape}')
        print(f'y shape: {y.shape}')
        return X, y
```

In [3]:

```
'''
Read training and test datasets
```

```
'''

Xdigitstrain, ydigitstrain =
readData('ZipDigits.train')

Ndigitstrain, pixels = Xdigitstrain.shape

assert(Ndigitstrain == ydigitstrain.shape[0])

Xdigitstest, ydigitstest = readData('ZipDigits.test')

Ndigitstest, pixels = Xdigitstest.shape

assert(Ndigitstest == ydigitstest.shape[0])


X shape: (7291, 256)
y shape: (7291, 1)
X shape: (2007, 256)
y shape: (2007, 1)
```

In [4]:

```
'''
Show images of handwritten digits
'''

def showKthImage(X, y, k):
    image = X[k, :].reshape((16, 16))
    plt.imshow(image, cmap='gray', vmin=-1, vmax=1)
    plt.title(f'First digit is a {y[k, 0]}')
    plt.xlim(0, 15)
    plt.ylim(15, 0)
    plt.xticks(range(16))
    plt.yticks(range(16))
    plt.tight_layout()
    plt.show()

showKthImage(Xdigitstrain, ydigitstrain, 0)
```



In [5]:

```
'''
Compute the augmented matrix with features
Helper Functions
'''

def computeIntensity(X):
    '''
    Input:
        X: a 2 dimensional N x 256 numpy array
            each row contains the values of 256 pixels
            from a 16 x 16 grayscale image of a handwritten digit
            each pixel has an intensity value between -1
and 1
    Output:
        intensities: a 2 dimensional N x 1 numpy array
            each row consists of a single
value representing the
            average pixel intensity of the
corresponding image
            See LFD Example 3.1
    '''
    print('computing intensity feature')
```

```

    N, d = X.shape
    print(f'Input shape {N}, {d}')
    '''

    TODO: Compute the intensity feature for N data
points
    '''

    print(f'Output shape {intensities.shape}')
    return intensities

def computeSymmetry(X):
    '''
    Input:
        X: a 2 dimensional N x 256 numpy array
            each row contains the values of 256 pixels
from a 16 x 16 grayscale image of a handwritten digit
            each pixel has an intensity value between -1
and 1
    Output:
        symmetries: a 2 dimensional N x 1 numpy array
            each row consists of a single
value representing the
                    "horizontal" symmetry of the 16 x
16 image about the vertical axis
                    See LFD Example 3.1
    '''
    print('computing symmetry feature')
    N, d = X.shape
    print(f'Input shape {N}, {d}')
    Ximgs = [X[n, :].reshape((16, 16)) for n in
range(N)]
    Ximgs_flipped = [np.flip(Ximgs[n], axis=1) for n in
range(N)]
    '''

    TODO: Compute the symmetry feature for N data
points
    '''

    symmetries = symmetries.reshape(N, 1)
    print(f'Output shape {symmetries.shape}')
    return symmetries

```

```

def computeAugmentedXWithFeatures(X):
    '''
    Input:
        X: a 2 dimensional N x 256 numpy array
            each row contains the values of 256 pixels
            from a 16 x 16 grayscale image of a handwritten digit
            each pixel has an intensity value between -1
            and 1
    Output:
        Xaug: a 2 dimensional N x 3 numpy array
            the augmented feature matrix
            the i-th row corresponds to the i-th row
            of X (and image represented by it)
            the 0-th column is the column of 1s
            the 1-st column is the column of average
            intensities
            the 2-nd column is the column of
            horizontal symmetries
    '''
    N, d = X.shape
    intensity = computeIntensity(X)
    symmetry = computeSymmetry(X)
    dummy = np.ones((N, 1))
    Xaug = np.concatenate((dummy, intensity, symmetry),
axis=1)
    # print(Xaug)
    print (f'Shape of augmented feature matrix:
{Xaug.shape}')
    return Xaug

```

In [6]:

```

'''
Compute the augmented matrix with features
'''

print('Computing augmented training feature matrix')

Xaugtrain = computeAugmentedXWithFeatures(Xdigitstrain)

```

```

Naugtrain, d = Xaugtrain.shape

print('Computing augmented test feature matrix')

Xaugtest = computeAugmentedXWithFeatures(Xdigitstest)

Naugtest, d = Xaugtest.shape

```

```

Computing augmented training feature matrix
computing intensity feature
Input shape 7291, 256
Output shape (7291, 1)
computing symmetry feature
Input shape 7291, 256
Output shape (7291, 1)
Shape of augmented feature matrix: (7291, 3)
Computing augmented test feature matrix
computing intensity feature
Input shape 2007, 256
Output shape (2007, 1)
computing symmetry feature
Input shape 2007, 256
Output shape (2007, 1)
Shape of augmented feature matrix: (2007, 3)

```

In [7]:

```

'''
Create the dataset with digits 1 and 5
'''

def indexDigits(y):
    '''
    Input:
        y: N x 1 2 dimensional numpy array; labels for
        handwritten digits
    Output:
        digit_idx: a dictionary; the keys are digits 0
-- 9
        for a digit k, digit_idx[k] is a

```

```

list identifying the rows labeled with digit k
'''
N = y.shape[0]
digit_idx = {}
for n in range(N):
    digit = ydigitstrain[n, 0]
    if not digit in digit_idx:
        digit_idx[digit] = []
    digit_idx[digit].append(n)
return digit_idx

```

In [8]:

```

'''
Construct the training and test sets for the rest of
the exercises on classifying 1s vs 5s
'''

digit_idx_train = indexDigits(Xaugtrain)
X1train = Xaugtrain[digit_idx_train[1], :]
N1train = X1train.shape[0]
print(f'number of 1s: {N1train}')
X5train = Xaugtrain[digit_idx_train[5], :]
N5train = X5train.shape[0]
print(f'number of 5s: {N5train}')
Xtrain =
Xaugtrain[digit_idx_train[1]+digit_idx_train[5], :]
ytrain = np.concatenate((np.ones((N1train, 1)),
-1*np.ones((N5train, 1))), axis=0)
Ntrain, d = Xtrain.shape
print(f'number of 1s and 5s: {Ntrain}')
print(f'Xtrain shape: {Xtrain.shape}, ytrain shape:
{ytrain.shape}')

digit_idx_test = indexDigits(Xaugtest)
Xtest =
Xaugtest[digit_idx_test[1]+digit_idx_test[5], :]
ytest =
np.concatenate((np.ones((len(digit_idx_test[1]), 1)),
-1*np.ones((len(digit_idx_test[5]), 1))), axis=0)
Ntest, d = Xtest.shape

```

```
print(f'number of 1s and 5s: {Ntest}')
```

```
print(f'Xtest shape: {Xtest.shape}, ytest shape: {ytest.shape}')
```

```
number of 1s: 1005
number of 5s: 556
number of 1s and 5s: 1561
Xtrain shape: (1561, 3), ytrain shape: (1561, 1)
number of 1s and 5s: 428
Xtest shape: (428, 3), ytest shape: (428, 1)
```

In [9]:


```
'''
Plot the training data
'''

fig, axs = plt.subplots(figsize=(5,5))
axs.scatter(X1train[:, 1], X1train[:, 2], marker='o',
color='blue', label='1')
axs.scatter(X5train[:, 1], X5train[:, 2], marker='x',
color='red', label = '5')
axs.set_xlabel('intensity')
axs.set_ylabel('symmetry')
axs.set_xlim(-1, 0.2)
axs.set_ylim(1, 2)
axs.set_title('1s and 5s')
axs.legend()
```

Out[9]:

```
<matplotlib.legend.Legend at 0x21998a93af0>
```





In [10]:

```
'''
Functions to compute the misclassification error
'''

def error(w, x, y, R=None):
    '''
    point-wise error measure for classification
    Input:
        w: a d x 1 2 dimensional numpy array
        x: a d x 1 2 dimensional numpy array
        y: a scalar value
        R: Risk weights; a dictionary
            whose keys are tuples (y, yhat) with
            value equal to the cost of predicting yhat
            when the label is y
    Output:
        error: misclassification error of hypothesis w
        on data point x with true label y
    '''
```

```

    '''
    TODO: compute the error made by hypothesis with
    weights w on data point x with label y
    '''

    return error

def E(w, X, y, R=None):
    '''
    point-wise error measure for classification
    Input:
        w: a d x 1 2 dimensional numpy array
        X: an N x d 2 dimensional numpy array
        y: an N x 1 2 dimensional numpy array
        R: Risk weights; a dictionary
            whose keys are tuples (y, yhat) with
            value equal to the cost of predicting yhat
            when the label is y
    Output:
        error: an N x 1 2 dimensional numpy array
            misclassification errors of hypothesis w
            on data points in X with true labels y
    '''

    # print(f'w shape {w.shape}, X shape {X.shape}, y
    shape {y.shape}')
    N = X.shape[0]
    '''
    TODO: compute the errors made by hypothesis with
    weights w on data points in X with true labels y
    '''

    return error

```

In [17]:

```

'''
Helper function to plot a linear separator
'''

def plotLinearSeparator(w, X, y, title=''):
    '''
    Plot data points a linear separator

```

```

Input:
    w: a d x 1 2 dimensional numpy array
    X: an N x d 2 dimensional numpy array
    y: an N x 1 2 dimensional numpy array
    title: a string
Output:
    error: misclassification error of hypothesis w
on data points in X with true labels y
'''
'''
Plot data points in X, y
'''
plusls = np.where(y == 1)[0]
minusls = np.where(y == -1)[0]
Xplusls = X[plusls, :]
Xminusls = X[minusls, :]
fig, axs = plt.subplots(figsize=(5,5))
axs.scatter(Xplusls[:, 1], Xplusls[:, 2],
marker='o', color='blue', label='1')
axs.scatter(Xminusls[:, 1], Xminusls[:, 2],
marker='x', color='red', label = '5')
axs.set_xlabel('intensity')
axs.set_ylabel('symmetry')
axs.set_xlim(-1, 0.2)
axs.set_ylim(1, 2)
'''
Plot separator
'''
pltxs = np.linspace(-1, 0.21)
pltys = - (w[0] + w[1] * pltxs) / w[2]
axs.plot(pltxs, pltys, color='green',
label='separator')
axs.set_title(title)
axs.legend()
plt.show()

```

In [18]:

```

'''
The Pocket algorithm (variant of the Perceptron
Learning Algorithm)
'''

```

```

def pocket(X, y, max_iters=1000, w_init=None):
    '''
    Implements the Pocket algorithm
    Input:
        X: A 2 dimensional N x d numpy array
            The i-th row X[i, :] contains features for
the i-th example in the training set
            X[i, 0] = 1
            X[i, 1], ... X[i, d] have values of features
        y: A 2 dimensional N x d numpy array
            y[i, 0] is the label associated with the i-
th example
        max_iters: an integer; maximum number of
iterations of PLA
        w_init: A 2 dimensional d x 1 numpy array
            intended to set initial weights for PLA
    Output:
        w_best: a d x 1 2 dimensional numpy array
            weights with lowest error on the input
training set X, y
    '''
    Eins = []
    ws = []

    Ein_best = np.infty
    w_best = 0

    w = np.zeros((d, 1))
    w = w + 0.0000001
    if not isinstance(w_init, type(None)):
        w = w_init

    for i in range(max_iters):
        Ein = E(w, X, y)
        Eins.append(Ein)
        ws.append(w)
        if Ein < Ein_best:
            Ein_best = Ein
            w_best = w

```

```

'''
    TODO: Complete this implementation of the
    Pocket algorithm
'''

plt.scatter(range(max_iters), Eins)
plt.xlabel('iteration')
plt.ylabel('in sample error')
plt.title('PLA')
plt.tight_layout()
plt.show()

print(f'Ein_best {Ein_best}, \nw_best \n{w_best}')

return w_best

```


In [20]:

```

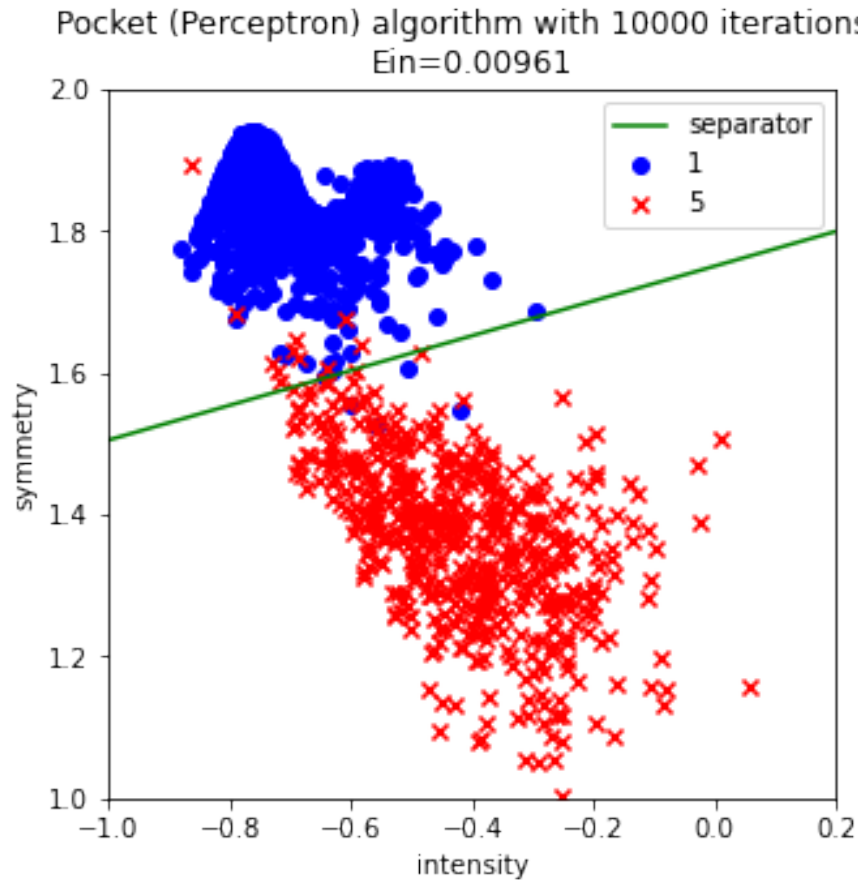
'''
Run the Pocket algorithm
'''

max_iters = 10000
w = pocket(Xtrain, ytrain, max_iters=max_iters)
Ein = np.round(E(w, Xtrain, ytrain),5)
plotLinearSeparator(w, Xtrain, ytrain, title=f'Pocket
(Perceptron) algorithm with {max_iters} iterations;
\nEin={Ein}')

```



```
Ein_best 0.009609224855861626,  
w_best  
[[-20.9999999 ]  
 [ -2.9436249 ]  
 [ 11.99648448]]
```



In [24]:

```
'''
The one-step optimal algorithm for Linear Regression
(See LFD Section 3.2.1)
'''
def linearRegression(X, y):
    '''
    Implements the one-step algorithm for Linear
    Regression (See LFD Section 3.2.1)
    Input:
        X: A 2 dimensional N x d numpy array
            The i-th row X[i, :] contains features for
            the i-th example in the training set
            X[i, 0] = 1
            X[i, 1], ... X[i, d] have values of features
        y: A 2 dimensional N x d numpy array
            y[i, 0] is the label associated with the i-
            th example
    Output:
```

```

        w: a d x 1 2 dimensional numpy array
        weights with lowest error on the input
training set X, y
'''
'''

    TODO: Implement the one-step optimal algorithm for
linear regression
'''

    print(f'Ein {Ein}, \nwlin \n{wlin}')
    return wlin

```

In [25]:

```

'''
Run Linear Regression followed by the Pocket algorithm
to classify 1s vs 5s
'''

wlin = linearRegression(Xtrain, ytrain)
Ein = np.round(E(wlin, Xtrain, ytrain), 5)
plotLinearSeparator(wlin, Xtrain, ytrain,
title=f'Linear Regression; \nEin={Ein}')

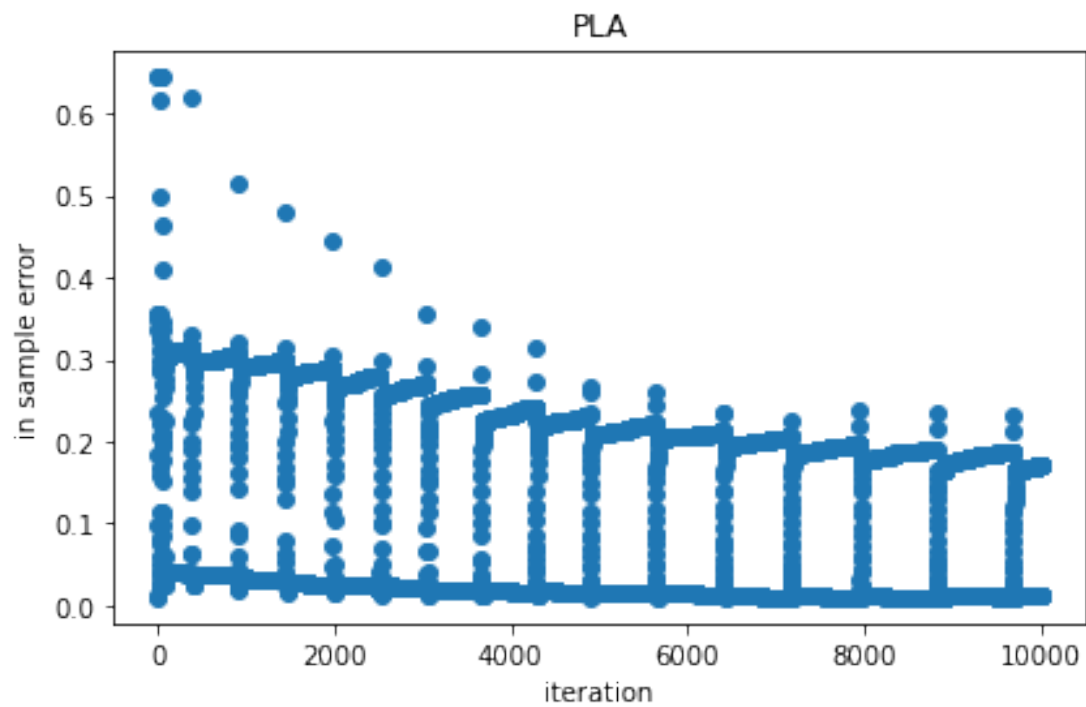

max_iters = 10000
w = pocket(Xtrain, ytrain, winit=wlin,
max_iters=max_iters)
Ein = np.round(E(w, Xtrain, ytrain), 5)
plotLinearSeparator(w, Xtrain, ytrain, \
                    title=f'Pocket algorithm after
Linear Regression with {max_iters} iterations;
\nEin={Ein}')
```

```

Ein 0.010249839846252402,
wlin
[[-6.05415811]
 [ 0.16071908]
 [ 3.85642483]]

```






```

Ein_best 0.009609224855861626,
w_best
[[-21.05415811]
 [ -2.77235905]
 [ 12.10354983]]

```

Pocket algorithm after Linear Regression with 10000 iterations;  
Ein=0.00961



In [26]:

```

'''
Gradient descent to minimize an arbitrary function
'''

def functionf(x, y):
    '''
    Computes the value of an arbitrary function in two
    variables at the input location
    '''
    '''

```

```

    TODO: Compute the value of the function at point x,
y
    '''
    return fval

def gradientf(x, y):
    '''
    Computes the gradient of an arbitrary function in
two variables at the input location
    '''
    '''
    TODO: Compute the gradient
    '''
    return df_by_dx, df_by_dy


def gradientDescent4f(x, y, eta=0.001, max_iters=100):
    '''
    Performs gradient descent to find the location at
which the value of an arbitrary function is minimized
    '''
    fvals = []
    for i in range(max_iters):
        # print(f'iteration {i}, x={x}, y={y}')
        fval = functionf(x, y)
        fvals.append(fval)
        grad = gradientf(x, y)
        '''
        TODO: Complete this implementation of gradient
descent for an arbitrary function f with two variables
        '''

        plt.plot(range(max_iters), fvals)
        plt.xlabel('iteration')
        plt.ylabel('value of f')
        plt.show()
    return x, y

max_iters = 1000
eta = 0.0001

```

```
gradientDescent4f(0.1, 0.1, eta=eta,  
max_iters=max_iters)
```



Out[26]:

```
(-0.24182360040042794, 4.698413037027736e-06)
```

In [27]:

```
'''  
The Logistic Regression algorithm with the cross  
entropy error measure  
'''  
  
def sigmoid(s):  
    '''  
    Implements the sigmoid function  
    '''  
    return 1 / (1 + np.exp(-1*s))  
  
def gradientCrossEntropyError(w, X, y):  
    '''  
    Computes the gradient of the cross entropy error  
    function on the dataset X, y at input weights w  
    See LFD Exercise 3.7  
    Input:
```

```

        w: a d x 1 2 dimensional numpy array
        X: an N x d 2 dimensional numpy array
        y: an N x 1 2 dimensional numpy array
    Output:
        gradient: a d x 1 2 dimensional numpy array
                   gradient of the cross entropy error
function on the dataset X, y at input weights w
'''
    N, d = X.shape
    grad = np.zeros((d, 1))
    for n in range(N):
        y_n = y[n, 0]
        x_n = X[n, :].reshape((d, 1))
        '''
        TODO: Complete this implementation to compute
the gradient at input weights w
'''

    return grad

def logisticRegression(X, y, eta=0.001, w_init=None,
max_iters=1000):
    '''
    Implements the gradient descent algorithm for
Logistic Regression
    See LFD Example 3.3
    Input:
        X: A 2 dimensional N x d numpy array
           The i-th row X[i, :] contains features for
the i-th example in the training set
           X[i, 0] = 1
           X[i, 1], ... X[i, d] have values of features
        y: A 2 dimensional N x d numpy array
           y[i, 0] is the label associated with the i-
th example
        eta: learning rate
        w_init: a d x 1 2 dimensional numpy array
                initial weights to start gradient
descent
        max_iters: maximum number of iterations of

```

```

gradient descent
    Output:
        w: a d x 1 2 dimensional numpy array
            weights with (approximately) lowest error on
the input training set X, y
'''

N, d = X.shape
w = np.zeros((d, 1))
if not isinstance(w_init, type(None)):
    w = w_init
Eins = []
for i in range(max_iters):
    Ein = E(w, X, y)
    Eins.append(Ein)
    grad = gradientCrossEntropyError(w, X, y)
    '''

    TODO: Complete this implementation of the
gradient descent algorithm for logistic regression
'''

plt.plot(range(max_iters), Eins)
return w

```


In [28]:

```

'''
Run the logistic regression algorithm to classify 1s vs
5s
'''

max_iters = 10000
eta = 0.5
w = logisticRegression(Xtrain, ytrain, eta=eta,
max_iters=max_iters)

```



In [29]:

```
plotLinearSeparator(w, Xtrain, ytrain, title=f'Logistic  
Regression with {max_iters} iterations; \nEin={Ein}')
```

